

Experiment 1

February 14, 2018

Simulation Details

Considered $K = 3$, $T = 5005$. Report statistics at $t = 1000, 3000, 5000$

The Bandit priors that were considered:

- Uniform: Draw the mean rewards for the arms from $[0.25, 0.75]$
- “HeavyTail”: We took the mean rewards to be randomly drawn from $\text{Beta}(\alpha = 0.6, \beta = 0.6)$. With this distribution it was likely to have arms that were at the extremes (close to 1 and close to 0) but also some of the arms with intermediate value means.
- Needle-in-haystack
 1. Medium - 9 arms with mean 0.50, 1 arm with mean 0.55 (+ 0.05)
 2. High - 9 arms with mean 0.50, 1 arm with mean 0.70 (+ 0.20)

Algorithms considered:

1. ThompsonSampling with priors of $\text{Beta}(1, 1)$ for every arm.
2. DynamicGreedy with priors of $\text{Beta}(1, 1)$ for every arm
3. Bayesian Dynamic ϵ -greedy with priors of $\text{Beta}(1, 1)$ for every arm and $\epsilon = 0.05$

Agent Algorithms considered:

1. HardMax
2. HardMaxWithRandom
3. SoftMax

Memory Sizes

1. 10
2. 25
3. 100

Simulation Procedure

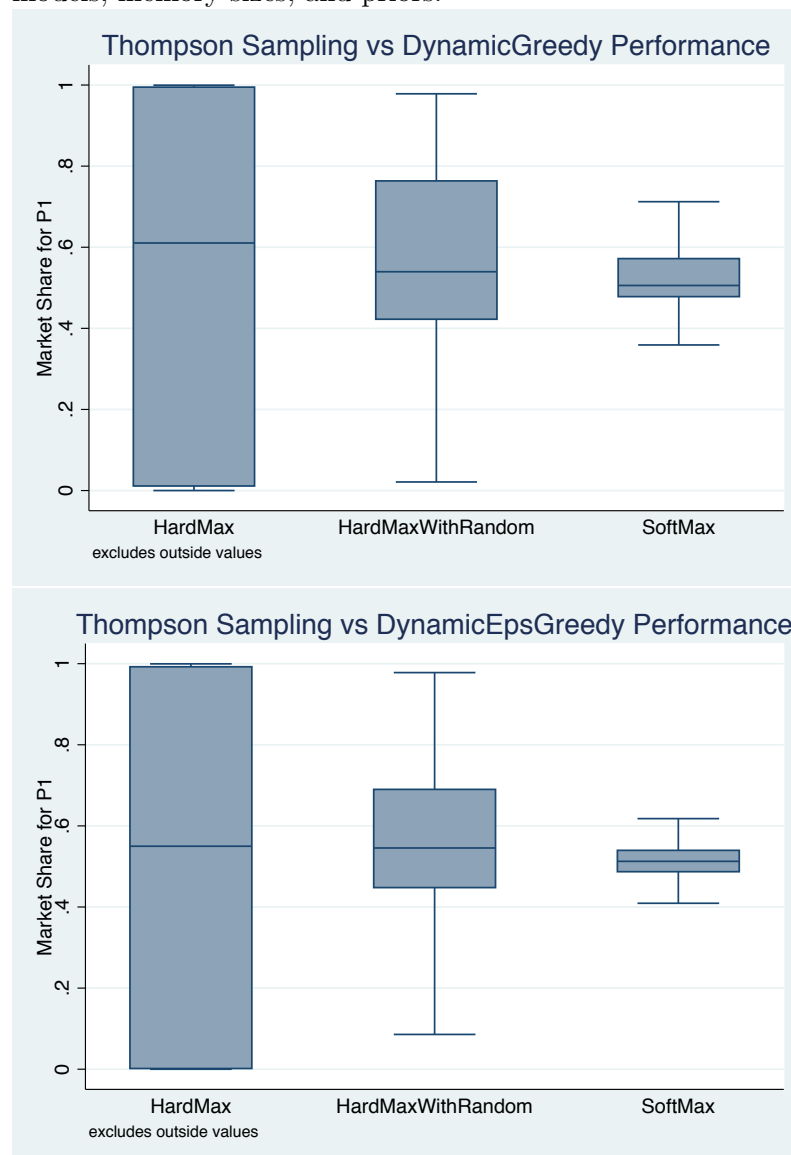
```
1: for Each prior  $p$  do
2:   for Each agent algorithm  $agent_{alg}$  do
3:     for Each principal algorithm pair  $principal_{alg1}, principal_{alg2}$  do
4:       for Each simulation  $i$  do
5:         Generate true distribution from  $p$  (except for needle-in- haystack, just use  $p$  itself)
6:         Run simulation for  $T$  periods
7:       end for
8:     end for
9:   end for
10: end for
```

Results

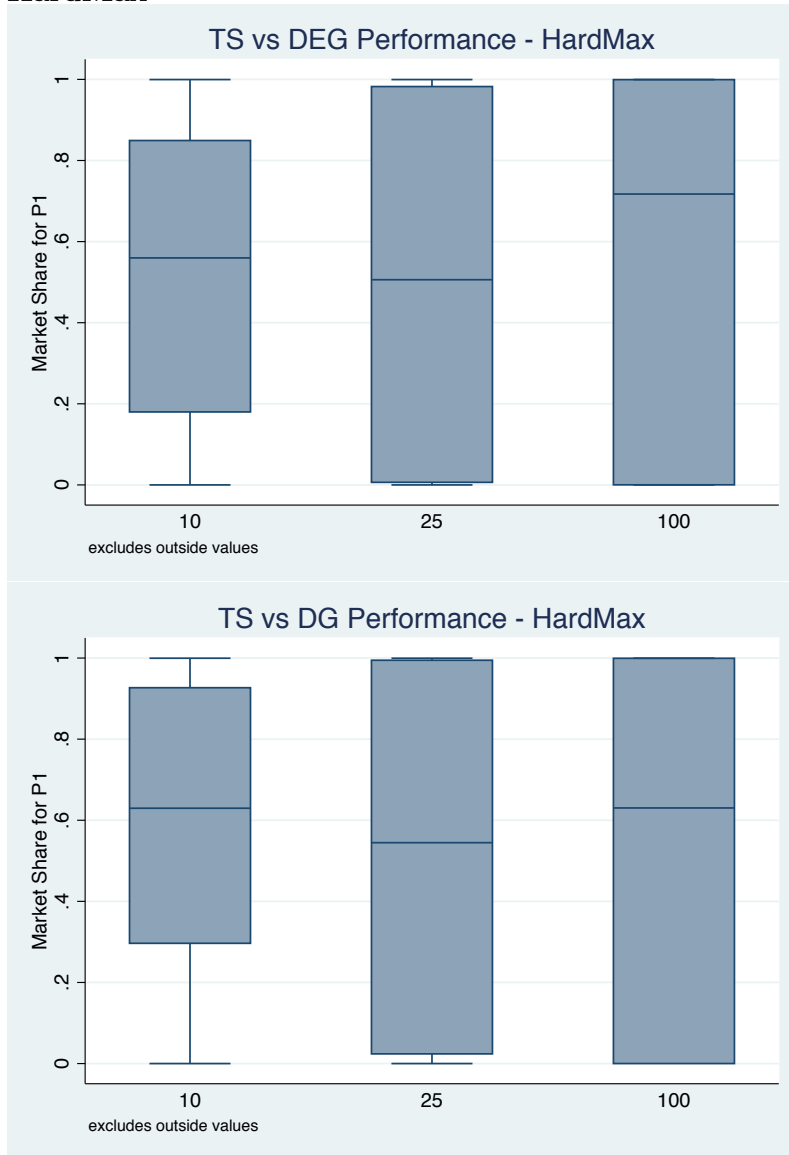
One thing which is ambiguous to define is the regret value to use when a principal never gets chosen in a given simulation. When calculating any of the aggregate regret statistics we drop these simulations, but we do record how many rounds have an undefined regret.

First, we'll restrict focus to $t = 5000$ and look at the performance of ThompsonSampling. Note that the y axis here represents the market share that the ThompsonSampling principal gets. Per-



formance of ThompsonSampling vs DynamicGreedy and DynamicEpsilonGreedy across all agent models, memory sizes, and priors:





Now, looking across different memory sizes for each agent model (still using each prior):
HardMax





HardMaxWithRandom



SoftMax

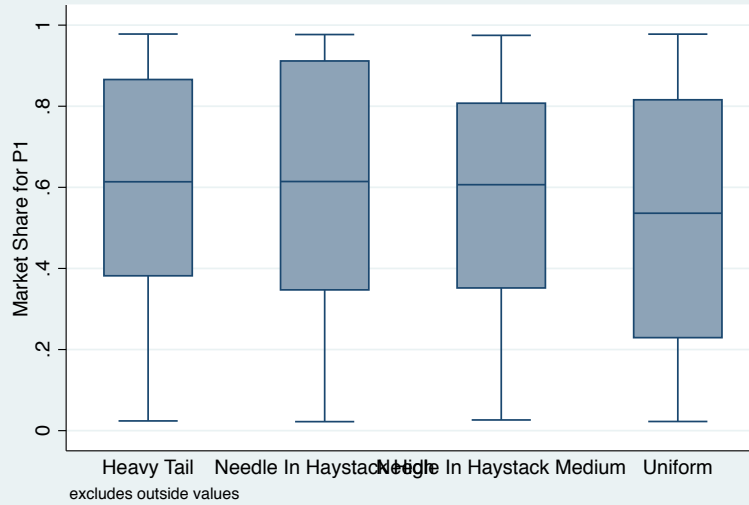


We'll now look fix memory size to be 100 and look at the performance across priors.
HardMax

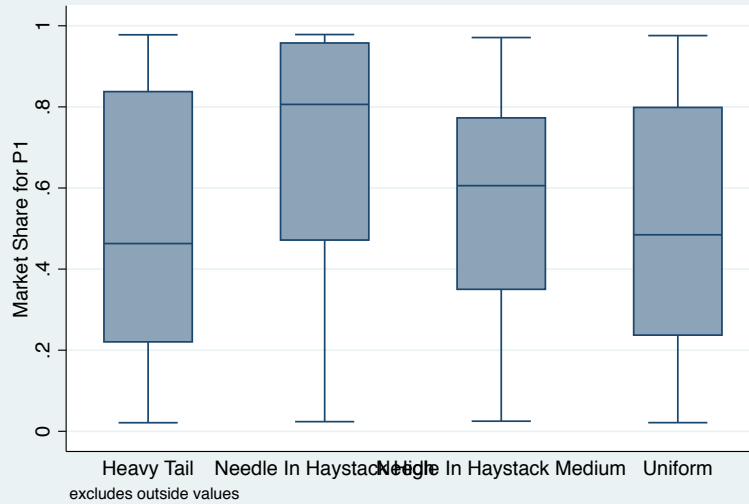


HardMaxWithRandom

TS vs DEG Performance - HardMaxWithRandom



TS vs DG Performance - HardMaxWithRandom



SoftMax

